

PA1 Report: BlisLab GEMM Optimization

Jingyu Wu, A16157847

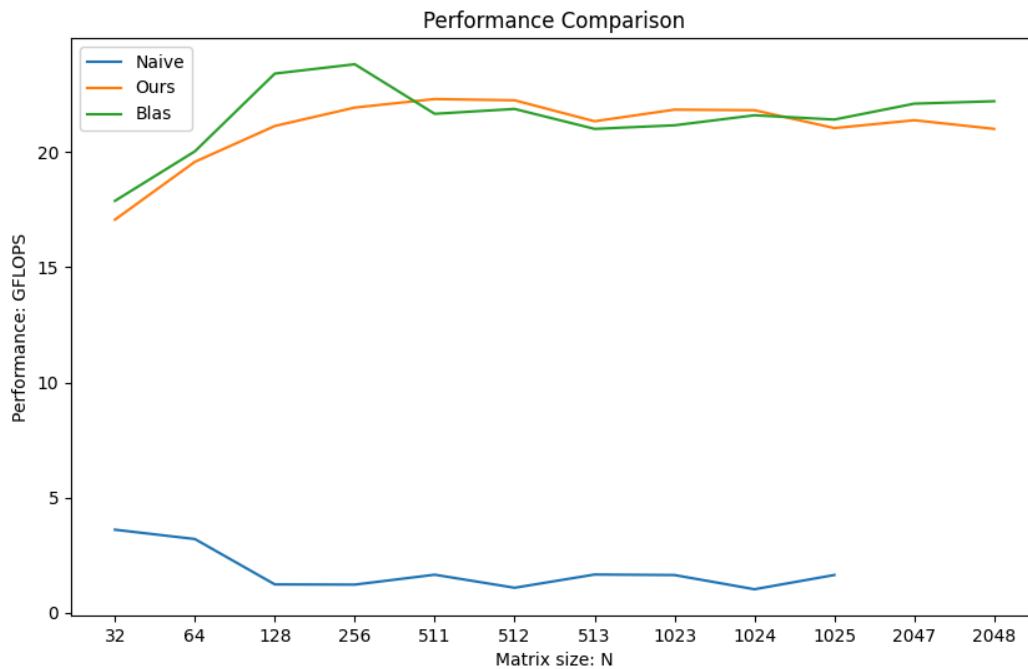
Haoxuan Wang, A16143053

Q1. Results

Q1.a. Optimized code performance:

N	Peak GF
32	17.07
64	19.58
128	21.135
256	21.935
511	22.305
512	22.25
513	21.34
1023	21.845
1024	21.82
1025	21.045
2047	21.385
2048	21.01

Q1.b. Performance Comparison (OpenBLAS, our code, naive code)



Note that the X-axis is equally spaced regardless of the N value.

Q2. Analysis

Q2.a. How does the program work

Key files:

benchmark.cpp: driver code to run matrix-matrix multiplications with different options and report results.

genDATA.sh: driver script (of the driver code) that takes average results over 20 runs on each matrix size.

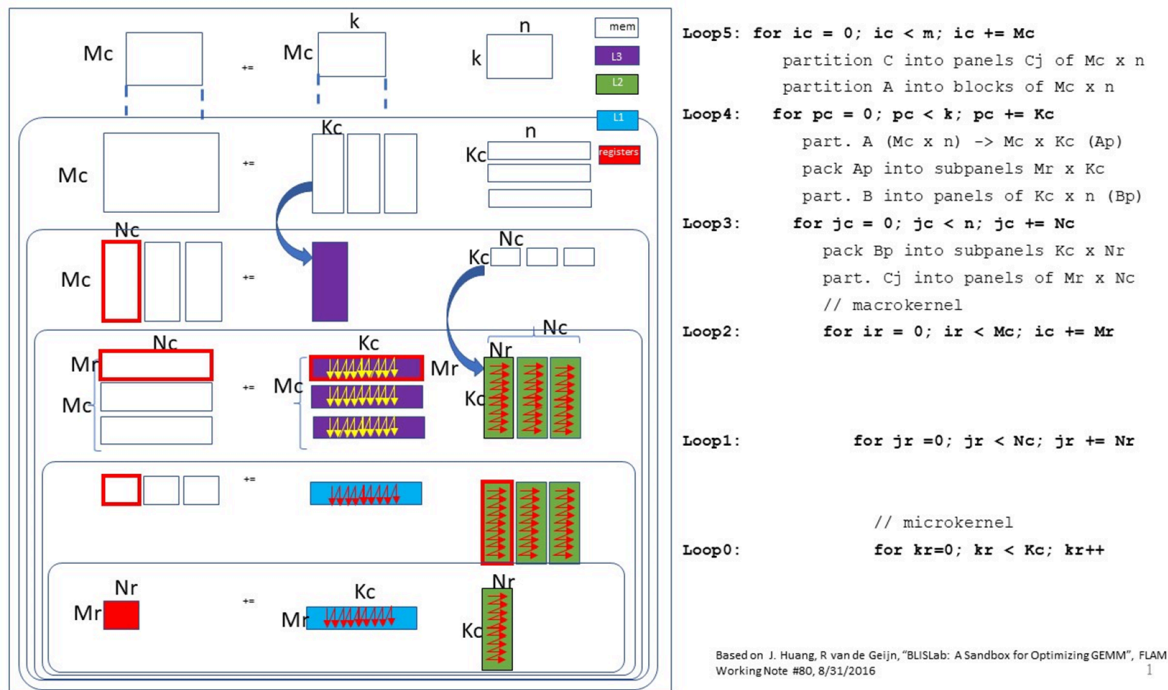
my_dgemm.c: main GEMM (general matrix multiplication) code, includes:

- Function `bl_dgemm()` for Loop 5 to 3, see below figure.
- Macro Kernel function `bl_macro_kernel()` for Loop 2 and 1.
- Packing functions `packA_mcxkc_d()` and `packB_kcxnc_d()`

bl_dgemm_ukr.c: highly optimized microkernel functions.

blislab/bl_config.h: definitions of hyperparameters, which include M_c , K_c , N_c , M_r , N_r in the below figure, and which microkernel to use.

Program Workflow:



The graph above shows how our program divides up the work of multiplying two big matrices into smaller blocks to enhance performance. Here is a further breakdown with regards to the code that better explains the blocking process:

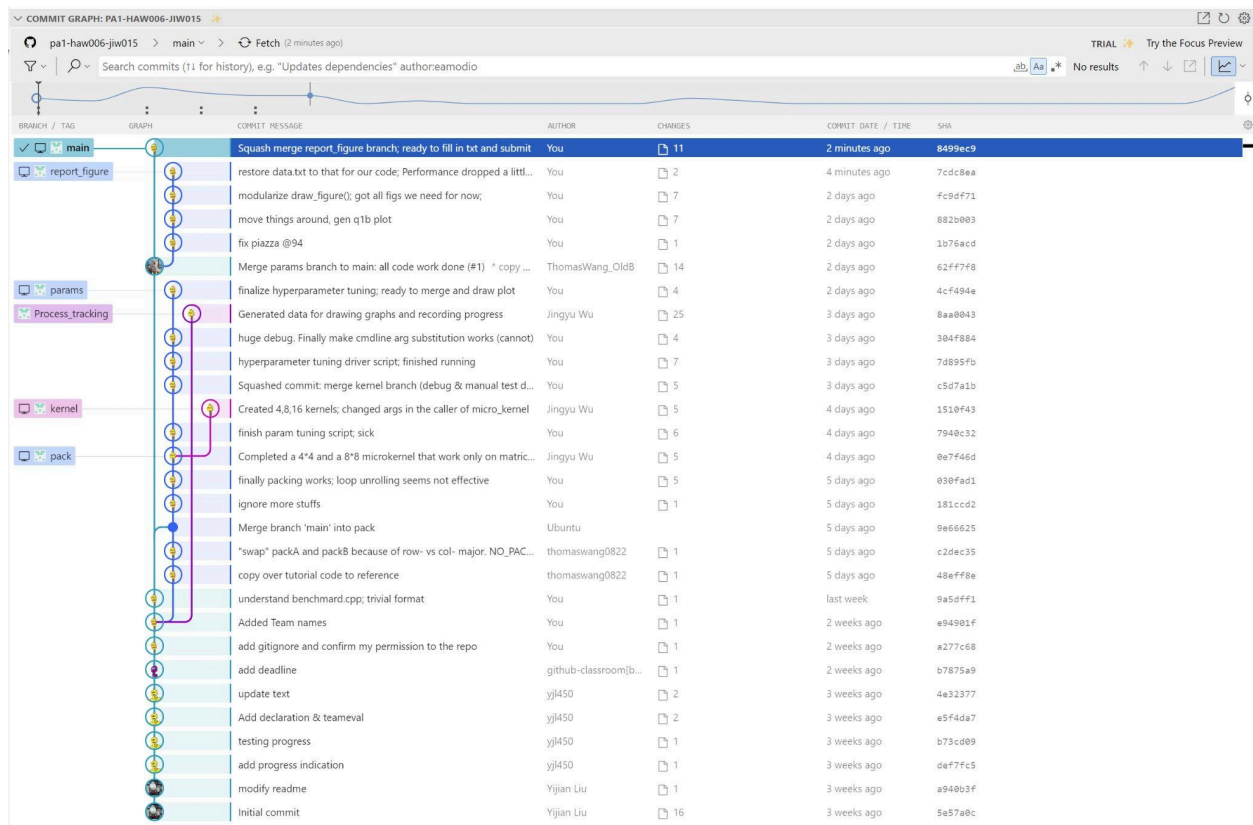
- 5th Loop (ic loop): Iterates over matrix A by blocks of rows (DGEMM_MC rows at a time). DGEMM_MC is a predefined size that specifies how many rows of A are considered in one block.
- 4th Loop (pc loop): Iterates over the shared dimension k (in $A[m*k] * B[k*n]$) taking blocks of size DGEMM_KC. This dimension dictates the depth of the multiplication, akin to the inner dimension in matrix multiplication.
 - Within the 4th loop, blocks of matrices A and B are prepared for efficient multiplication:
 - Packing A: Blocks of matrix A are packed using `packA_mcxkc_d`, ensuring they are stored in column-major order within the packed block. This

function also handles padding with zeros if the block does not fully match DGEMM_MR.

- Packing B: Similarly, blocks of matrix B are packed using packB_kcxnc_d, but in row-major order. Padding with zeros is similarly managed for unmatched DGEMM_NR.
- 3rd Loop (jc loop): Iterates over matrix B by blocks of columns (DGEMM_NC columns at a time).
- Nested Inner Loops: These loops set the stage for the actual multiplication:
 - They determine the exact size of each block considering edge cases where the dimensions of the matrices do not align perfectly with the block sizes (DGEMM_MC, DGEMM_NC, DGEMM_KC).
 - The micro-kernel, referenced as bl_micro_kernel, is invoked for each block pair from matrices A and B. The results are accumulated directly into the appropriate block of matrix C.

Q2.b. Development process:

We ran the benchmark program after each step, but we didn't record the data or generate the plot until we finished everything. Nevertheless, we maintained good code structure and commit histories such that we could easily restore the code to any previous state to draw plots and compare performance. This GitLens commit graph helps with tracking progress:



1. Starter Code:


We first tried running the starter code, which used the most basic ijk micro-kernel.

2. Packing:

The first job is the packing routines. After that, the performance reached just above 2.0 GFLOPS, which is below our expectations. This suggests that packing overhead counters most of the acceleration it brings.

3. Loop Unrolling:


We then added the loop unrolling in the packing code with some easy changes. There was an improvement, but not significant at all.



In the plot legend, “+=” means the optimization technique is built upon previous ones in the order above.

4. First microkernel


We started to build microkernels. The first one operates 4 numbers at a time in a vector. The new kernel immediately boosted the performance from under 3 GFLOPS to around 18 GFLOPS.



In this graph, the data of the blue line ('bl_dgemm_ukr') are the same as those of '+Unrolling' in the previous plot.

5. Write more microkernels and do manual testing

We then implemented more microkernels with the same style but operated on more numbers (8 and 16) at a time. The improvement isn't as significant, but still, they push our performance above 20 GFLOPS. Here is their comparison:



6. Automated parameter tuning

To begin with, we figured out what values of the hyperparameters (namely, Nr, Mr, Kc, Nc, Mc) should be taken by examining cache sizes with `$ lscpu | grep -i cache`. Some brief derivation is noted down in “results/cache_hyperparm.txt”.

Writing the Python script that automates the process of “compile with different macros (hyperparameters), run, record result” isn’t easy. In particular, the biggest difficulty comes from passing key-value pairs to substitute the preprocessor macros defined in the code. Hours were spent before we realized that to use ‘-D’ to pass in key-value pairs, those macros in the code must be commented out because they take high priority.

After all, we decided that (Nr, Mr, Kc, Nc, Mc) = (4, 16, 256, 128, 1024) gives the best result in all combinations. And we use them to generate data and draw the plot in Q1.b.

Q2.c. Irregularities in the data

N	Performance (GFlops)
32	17.07
64	19.58

128	21.135
256	21.935
511	22.305
512	22.25
513	21.34
1023	21.845
1024	21.82
1025	21.045

Based on the table above and the graph in section Q1.b, we found out that the performance of our matrix multiplication subroutine reaches its peak at around $n=512$, where it no longer improves when n gets larger.

We also noticed that the performance metrics are higher at values like $n=511$ or $n=512$ compared with $n=513$. We conjecture that it may be due to the additional overhead created by the padding in the packing functions.

Q2.d. Supporting data

Cache Sizes and hyperparameter choices:

Using the command `"lscpu | grep -i cache"`, we found out the sizes for the three levels of caches in the machine:

L1d cache: 64 KiB (1 instance)

L1i cache: 64 KiB (1 instance)

L2 cache: 1 MiB (1 instance)

L3 cache: 32 MiB (1 instance)

Hence, we can decide the upper bounds of our hyperparameters (N_r , M_r , K_c , N_c , M_c) based on the following calculations:

N_r : has to be 4 for a row of C to fit into one vector register.

Mr: Based on our implementation of micro-kernels, $M_r = 4, 8, \text{ or } 16$

Kc: $K_c * M_r$ block of A has to fit into the L1 cache (64KB), so $K_c * M_r < 64 * 1024 / 8 = 8192$

Nc: $N_c * K_c$ block of B has to fit into the L2 cache (1 MB), so $N_c * K_c < 1024^2 / 8 = 131072$

Mc: $M_c * K_c$ block of A has to fit into the L3 cache (32 MB), so $N_c * K_c < 32 * 1024^2 / 8 = 4194304$

Difference from the BLISlab tutorial

The first modification involves changing to row-major matrices for the input matrices A and B and for the output matrix C, differing from the BLISLab tutorial which uses column-major order. This adjustment is necessary as our machine uses row-major order, and we also need to adapt the skeleton code to accommodate this.

The second modification involves managing specific scenarios by padding the matrices with 0s, particularly in cases where N is not evenly divisible by the various block sizes. However, paddings introduce additional overhead, which affects the performance when N is small.

Q2.e. Future work

Our PA, or BLAS (Basic Linear Algebra Subprograms) in general, is concerned with optimizations in low-level CPU routines. After this, we could look at other sources of acceleration, such as multi-core, distributed (multi-device) computations. Also, in this AI era, multiplications of huge matrices arise from the training of gigantic neural networks. Therefore, storage is also an important issue.

Our performance could likely be improved if we pick different hyperparameters and microkernels for different matrix sizes. The code shouldn't be difficult to write, but determining the condition robustly needs a lot more benchmarking.

Q3. References

Class resources:

PA 1 Instructions, Lecture slides, Discussion slides

Papers:

[1]Jianyu Huang and Robert A Van de Geijn. 2016. BLISlab: A sandbox for optimizing GEMM. arXiv preprint arXiv:1609.00076 (2016).

[2]Kazushige Goto and Robert A van de Geijn. 2008. Anatomy of high-performance matrix multiplication. ACM Transactions on Mathematical Software (TOMS) 34, 3 (2008), 1–25.

[3]Field G Van Zee and Robert A Van De Geijn. 2015. BLIS: A framework for rapidly instantiating BLAS functionality. ACM Transactions on Mathematical Software (TOMS) 41, 3 (2015), 1–33.

Links (to look for some easy commands on StackOverflow, and some other tools):

<https://stackoverflow.com/a/30211477/14697376>

<https://stackoverflow.com/a/9890599/14697376>

<https://stackoverflow.com/a/37708190/14697376>

BibTex Converter: <https://asougi.github.io/bibtex-converter/>

Matplotlib API: <https://matplotlib.org/stable/index.html>

Q4. Extra Credit

Unfortunately we didn't have time to complete any extra credit option.